



(a) Expert activation frequency



(b) Per-GPU compute load



(c) NIC-to-NIC dispatch traffic